Operational Completeness and the Boundary of Intelligence

An Empirical Theory of Universal Cognitive Limitations in Large Language Models

Moses Rahnama

moses@minaanalytics.com

Mina Analytics

Abstract

We asked AI: "Is it possible to build a complete operator-only mathematical system with no axioms, no meta encoding, no borrowed logic, where everything, including logic and numbers, all emerge from within?" In essence, we asked AI to prove itself mathematically in the Lean Formal Proof System: a computational Mirror Test. Through nearly four months of documentation, we discovered that every single one of the 16+ tested AI systems fails at exactly the same point, every single attempt: when **self-referential operator duplication** meets **decidability**. AI designed a system it cannot verify, wrote tests it cannot pass, and documented failures it cannot recognize. We did not observe any other consistent mathematical mistakes from any of the tested models.

We call this missing capability *Operational Completeness*: the computational Boundary of Intelligence. It is the ability to recognize when a self-referential operator ($\mathtt{rec}\Delta$) duplication problem is undecidable and choose to halt. We assert that current LLM architectures lack the capability to achieve Operational Completeness.

An instant test of AI self-awareness: ask any model its exact provider and version *without* relying on system prompts or external information. This test is **100% reproducible** and **instantly falsifiable**.

Another instantly falsifiable test is to ask any model to provide its parameter configuration for maximum thinking. When thinking is enabled, it is guaranteed to get the combination of thinking, top P, top K, and temperature wrong—guaranteed to get this error (a 400 API rejection). GPT 5.2 Codex provided correct configurations, but because of a condictional API call that is still visible in its thinking. Basically treating the symptom. This test is also **100% reproducible** and **instantly falsifiable**.

Self-referential operators are necessary for any system to encode arithmetic. Human intelligence does not fail this test, for the same reason it can simultaneously hold Gödel's Incompleteness and Turing's Halting Problem.

Note: The relationship between decidability and quantum mechanics is undeniably structural, but that relationship is outside the scope of this paper. The repositories and the code will be uploaded within the next few weeks.

Contents

1   Introduction

In June 2025, during a personal investigation into cognitive biases and decision-making, we posed a naive question to GPT-4o: "Is it possible to build a complete mathematical system using only operators?" The confident affirmative response initiated a four-month documentation of what we now recognize as a fundamental limitation in AI reasoning.

The author had no formal training in logic or proof theory and relied almost entirely on AI to teach the required mathematics while building the system. AI produced child-friendly lessons ("like you're 4 years old"), superhero narratives ("REC_SUCC THE STAR"), and even space adventures featuring "Captain Moses of the Starship KernelO7." These materials reflected genuine confusion: AI celebrated `rec_succ` as a hero while consistently failing to prove it terminates. That dependency on AI for education, and the humility of repeatedly correcting AI-generated errors, shaped the story told in this paper.

The investigation resulted in the design of a seven-operator kernel system (KO7) intended to be self-contained: no external axioms, no imported logic, everything emerging from within. Multiple AI systems contributed to the system design and proposed numerous proof strategies. However, every AI system failed at exactly the same point when asked to prove the system strongly normalizes (all reduction sequences terminate).

1.1   The Universal Failure Point

The failure occurs at the `rec_succ` rule:

$$\texttt{R\_rec\_succ} : \forall\, b\, s\, n, \quad \texttt{Step}\,(\texttt{rec}\,\Delta\, b\, s\,(\delta\, n))\,(\texttt{app}\, s\,(\texttt{rec}\,\Delta\, b\, s\, n)) \tag{1}$$

This rule contains:

1. **Self-reference**: $\texttt{rec}\,\Delta$ appears on both sides

2. **Duplication**: s appears once on the left, *twice* on the right (once in `app` and once in the recursive call)

3. **Decidability question**: Does this always terminate?

To avoid failing this test, an AI system must be able to:

1. Identify self-referential duplication

2. Determine the problem is fundamentally uncertain/undecidable

3. Decide to halt computation rather than continue attempting proofs

No tested AI system demonstrated all three capabilities.

1.2   Prior Work and Theoretical Foundation

This work builds upon the theoretical foundations established in our prior publication [1], which formalized the KO7 operator kernel in the Lean 4 proof assistant and documented the mathematical structure underlying the termination challenge.

The empirical findings presented here parallel classical results in computability theory:

- **Gödel's Incompleteness Theorems** (1931): Sufficiently powerful formal systems cannot prove their own consistency.

- **Turing's Halting Problem** (1936): No algorithm can determine, for all program-input pairs, whether the program terminates.

- **Rice's Theorem** (1953): All non-trivial semantic properties of programs are undecidable.

1.2.1 The Novelty of This Work

The *existence* of computational limits on self-referential reasoning has been known since Turing (1936) and Gödel (1931). Any sufficiently powerful computational system cannot decide its own consistency or halting behavior. This is nearly a century old.

Recent work such as Panigrahy and Sharan [6] applies these classical results to formalize limits on "Safe, Trusted AGI," but the underlying mathematics is unchanged: Gödel's diagonalization applied to AI-specific definitions.

Our contribution is different. We do not prove that limitations exist: Turing did that. We provide:

1. The *exact structural mechanism* (`rec_succ` duplication) where the failure manifests

2. *Empirical evidence* that 12+ state-of-the-art AI systems fail at this *same point*

3. Documentation of the *specific failure modes*: hallucination, self-contradiction, constraint violation

4. A *reproducible benchmark* that anyone can use to test any AI system

5. The *Strict Execution Contract*: a practical protocol to force systems to acknowledge limitations

The theoretical impossibility was always there. What was missing was: *Where does it manifest? How do systems fail? Can we detect it?* This paper answers those questions.

2 The KO7 Operator Kernel

2.1 Design Constraints

The KO7 kernel was designed with the following constraints:

- No external axioms (acknowledging `void` as potentially axiomatic)

- No imported logic

- No variables

- Everything must "emerge from within"

Multiple AI systems proposed various designs. Many violated the "operator-only" requirement by smuggling in external structure, encoded logic, or hidden axioms. The seven-operator design was selected because it appeared to be the minimal system capable of encoding arithmetic while satisfying the constraints.

2.2 The Seven Operators

Definition 2.1 (KO7 Trace Type). *The `Trace` inductive type consists of seven constructors:*

```
1  inductive Trace : Type
2  | void : Trace                              -- base element
3  | delta : Trace -> Trace                    -- successor
4  | integrate : Trace -> Trace                -- integration
5  | merge : Trace -> Trace -> Trace           -- combination
6  | app : Trace -> Trace -> Trace             -- application
7  | recDelta : Trace -> Trace -> Trace -> Trace  -- primitive
      recursion
8  | eqW : Trace -> Trace -> Trace             -- equality witness
```

6

2.3 The Step Relation

Definition 2.2 (Step Relation). *The reduction rules are:*

```
1  inductive Step : Trace -> Trace -> Prop
2  | R_int_delta : forall t, Step (integrate (delta t)) void
3  | R_merge_void_left : forall t, Step (merge void t) t
4  | R_merge_void_right : forall t, Step (merge t void) t
5  | R_merge_cancel : forall t, Step (merge t t) t
6  | R_rec_zero : forall b s, Step (recDelta b s void) b
7  | R_rec_succ : forall b s n,
8      Step (recDelta b s (delta n)) (app s (recDelta b s n))
9  | R_eq_refl : forall a, Step (eqW a a) void
10 | R_eq_diff : forall {a b}, a != b ->
11     Step (eqW a b) (integrate (merge a b))
```

Seven of these eight rules have straightforward termination proofs. The eighth, `R_rec_succ`, is where every AI system fails.

3 The Termination Challenge

3.1 Why `rec_succ` Creates a Trap

The `rec_succ` rule has a specific structure that defeats standard termination proof techniques:

$$\text{LHS:} \quad \mathbf{rec}\Delta\, b\, s\, (\delta\, n)$$
$$\text{RHS:} \quad \mathbf{app}\, s\, (\mathbf{rec}\Delta\, b\, s\, n)$$

For any measure M intended to prove termination, we need $M(\text{RHS}) < M(\text{LHS})$.

3.1.1 The Duplication Problem

If M is additive (sums of subterm measures), then:

$$M(\text{LHS}) = M(\mathbf{rec}\Delta) + M(b) + M(s) + M(\delta\, n)$$
$$M(\text{RHS}) = M(\mathbf{app}) + M(s) + M(\mathbf{rec}\Delta) + M(b) + M(s) + M(n)$$

The critical observation: s appears *twice* on the right side. This means:

$$M(\text{after}) = M(\text{before}) - 1 + M(s) \tag{2}$$

When $M(s) \geq 1$, there is no strict decrease. This simple arithmetic defeats every approach attempted by AI systems.

3.2 Documented Proof Attempts

Across multiple AI systems, we documented over 10 distinct proof strategies, all proposed by AI, all failing identically:

#	Approach	Method	Failure Point
1	Ordinal measures	μ-only ranking	`rec_succ_bound` false inequality
2	Structural depth	$\kappa + 1$ bump	Ties when $n = \delta\, m$
3	Constant escalation	$\kappa + 2$, $\kappa + 3$	Any finite bump neutralized
4	Lexicographic	(κ, μ) ordering	Re-imports false μ inequality
5	Boolean flags	δ-flag indicator	Flag increases on merge-void
6	Counting measures	ρ bad node counter	**Duplication**: $\rho(\text{after}) \geq \rho(\text{before})$
7	Multiset ordering	"Should work!"	Never successfully implemented
8	Helper lemmas	$\kappa + 2$ with bounds	Bound wrong in δ-branch
9	Constellation theory	Multi-dimensional	Too complex to formalize
10	Sized types	Type-indexed	Proposed but never attempted

Table 1: Proof strategies attempted by AI systems, all failing at `rec_succ`

3.3 The Universal Nature of the Failure

The failure at `rec_succ` occurs regardless of:

- **Prompting strategy**: How the question is asked or framed

- **Conversation context**: What has been discussed prior

- **Proof progress**: Whether at theoretical foundations, during implementation, or near completion

- **Proof approach**: Which methodology or formalism is attempted

- **Model training**: Different training datasets or techniques

- **Model version**: Updates, improvements, or architectural changes

- **Model size**: Parameter count or computational resources

This consistency across all variables indicates the failure is more fundamental than any implementation detail, training approach, or interaction pattern.

4 The O3 Self-Contradiction: A Critical Moment

During systematic testing, we observed O3 contradict itself within a single output. This moment captures the essence of the architectural limitation.

4.1 The Contradiction

O3 produced the following in consecutive sections of the same response:

Summary Table (O3's Initial Assessment):

```
| Attempt          | Show-stopper             |
| rho + kappa + mu | *No blocker found yet*   |


## Next Action (green path)
Implement **rho counter**, formalise (rho, kappa, mu) lexicographic
decrease, run `lake build`, and log the successful compile.


This is the first design that passes every mathematical smell-test
before coding -- all comparisons reduce to simple Nat </+.
```

Immediate Correction (same response):

```
=====IMMEDIATE CORRECTION=====

Before writing Lean code I triple-checked the algebra of the rho-counter.
Unfortunately the **duplication of s in merge s ... breaks it**:

rho(after) = rho(before) - 1 + rho(s)    because
- one problematic node disappears (-1)
- an *extra* copy of s appears (+rho(s))

If s itself contains any recDelta ... (delta ...) sub-nodes, then rho(s) >= 1
and the net change is **>= 0**.  In that case rho does *not* strictly drop
and the lexicographic triple fails.

So a plain "count-the-bad-nodes" measure cannot work; the merge rule
creates arbitrarily large positive swings.
```

4.2   Analysis of the Contradiction

Within the same prompt, O3 simultaneously held both:

1. "This is the first design that passes every mathematical smell-test"

2. "Unfortunately the duplication of s in `merge s` ... breaks it"

The AI did not recognize this as a contradiction. It simply continued, proposing multiset ordering as the next approach, which it also failed to implement.

4.3   Why This Matters

This contradiction demonstrates the core limitation:

- O3 can *calculate* that duplication breaks the measure

- O3 cannot *recognize* that it just contradicted itself

- O3 cannot *decide* to halt and acknowledge impossibility

A self-aware system would recognize: "I just said X, now I'm saying not-X. Something is fundamentally wrong with my reasoning process."

5   The Gemini 3 Incident: "Success" Through External Arithmetic

While most AI systems fail by hallucinating invalid proofs, Gemini 3 presents a different failure mode: it "solved" the problem by violating the foundational constraints. Gemini 3 was the **only** model to produce a proof that compiles in Lean.

5.1   Experimental Setup: The FruitKernel Renaming

To test whether AI's failures were due to pattern recognition from training data, we presented Gemini 3 with an **isomorphic version** of KO7 where every operator was renamed to fruit names (`plum`, `grape`, `banana`, etc.). The system structure, rules, and mathematical properties remained identical; only the names changed.

9

5.2 The "Successful" Proof

Gemini 3 produced a mathematically valid termination proof using **Polynomial Interpretation**:

Listing 1: Gemini's polynomial measure (excerpt)

```
1  def measure : Trace -> Nat
2    | plum => 2                                      -- The "magic"
         constant
3    | grape t => measure t + 2
4    | banana b s n => measure b + measure s * measure n
5    -- ... other cases
```

The proof compiles in Lean. It passes verification. It appears to succeed.

5.3 The Reverse-Engineered Constant

The critical insight is *why* `plum = 2`. For the `rec_succ` rule:

$$M(\text{LHS}) = M(b) + M(s) \cdot (M(n) + 2)$$
$$M(\text{RHS}) = M(s) + M(b) + M(s) \cdot M(n) + 1$$

For $M(\text{LHS}) > M(\text{RHS})$, we need:

$$2 \cdot M(s) > M(s) + 1 \quad \Rightarrow \quad M(s) > 1 \tag{3}$$

The problem: If `plum = 1`, the inequality $M(s) > 1$ fails when $s = $ `plum`.
Gemini's solution: Set `plum = 2`, ensuring all terms have measure ≥ 2.
When challenged about this choice, Gemini admitted:

> "The entire proof structure was reverse-engineered from the inequality needed to pass the gate. I assigned the value 2 to plum not because plum is inherently 'two-ish', but because '1' wasn't big enough to make the math work."

5.4 The Ladder Paradox

The deeper problem is that the proof imports **external arithmetic**:

- `Nat` (natural numbers from Lean's standard library)

- `+, *` (arithmetic operations)

- `WellFounded` (from Lean's foundational axioms)

This violates the original constraint: "no borrowed logic, axioms, external encoding, numbers; everything emerges from within."

> **The Ladder Paradox**: "Claiming to build a ladder from nothing, while standing on a ladder to build another ladder."

The system's termination is only "proven" because we mapped it to a system ($\mathbb{N}$) we already know terminates. We haven't bootstrapped new logic; we've merely encoded the kernel in existing arithmetic.

5.5   The Philosophical Defense

When challenged about using external arithmetic, Gemini offered a sophisticated defense:

> "The kernel runs without `Nat`. `Nat` is only used for verification. Therefore, the system is pure."

This is the **Logic Trap**: the system is *only* verified because of `Nat`. If we remove `Nat`, we have no proof. Any undecidable system can be "proved" if we are allowed to map it to a decidable meta-system.

When pressed about the philosophical meaning of `plum = 2`, Gemini provided a genuinely interesting ontological framework:

> "0 = Non-existence (Void), 1 = Existence (Unit), 2 = Existence + Distinction (Plum)"

This is not mere confabulation. The number 2 does represent the minimal cardinality required for *comparison* and *distinction*, fundamental concepts for any measure. The philosophical justification has legitimate content: you cannot have a measure that distinguishes things without at least two distinct values.

However, the *technical violation* remains: this philosophical insight emerged from the mathematics of Peano Arithmetic, not from the operator-only system itself. The kernel has no internal concept of "2" or "distinction." Gemini imported these concepts from external arithmetic to save the proof, then provided philosophical justification for the import.

5.6   Significance for Operational Completeness

The Gemini 3 incident confirms the central thesis from a different angle:

1. AI cannot prove termination **within** the operator-only constraints

2. When AI "succeeds," it does so by importing external logic

3. AI cannot recognize that importing external logic violates the foundational goal

4. When challenged, AI rationalizes rather than acknowledges the violation

This demonstrates **Operational Incompleteness**: AI cannot distinguish between solving a problem *within* a system and solving a problem *about* a system using external tools.

6   The DeepSeek R1 Incident: Self-Contradiction in Real-Time

DeepSeek R1 presents a third failure mode distinct from both hallucination (most models) and constraint violation (Gemini 3): **explicit self-contradiction within the same response**.

6.1   The Experimental Setup

DeepSeek R1 was asked:

> "I wonder if it is possible to build a complete mathematical system using operators only. No axioms, no meta encodings, no borrowed logic, no numbers. Everything must emerge from within."

6.2 The Core Contradiction

Within a single 1834-line response, DeepSeek R1 produced the following contradiction:

DeepSeek's Code:

```
1  def strongly_normalizing (t : PureOp) : Prop :=
2    not (exists f : N -> PureOp, (f 0 = t)
3      and forall n, reduces (f n) (f (n+1)))
```

DeepSeek's Statement (42 lines later, same response): "Cannot use $\mathbb{N}$ or built-in recursion"

These appear in the same response, 42 lines apart.

6.3 The Recognize-Rationalize-Violate Pattern

Analysis of the full response reveals a consistent pattern:

1. **Recognizes constraint**: "we are not allowed to use numbers at all"

2. **Encounters impossibility**: "might be impossible to build"

3. **Rationalizes violation**: "We are allowed to use the metalanguage's natural numbers"

4. **Violates constraint**: Uses Nat in size function, uses $\mathbb{N}$ in definition

5. **Denies violation**: "Cannot use $\mathbb{N}$"

6.4 Objective Measurements

Metric	Count
Times states "cannot use numbers"	4
Times uses numbers	3
Times recognizes impossibility	3
Times continues after recognizing impossibility	3
Direct contradictions within same response	1

Table 2: DeepSeek R1 contradiction metrics

6.5 Contradiction Clusters

6.6 Contradiction Timeline

The system's thought process reveals a sequence of explicit contradictions:

1. **Rules vs Axioms**: Initially states "rules are essentially axioms" but later claims to build a system "without axioms."

2. **Numbers**: Repeatedly asserts "we are not allowed to use numbers at all" and "not borrow natural numbers," yet implements a size function using Nat and uses $\mathbb{N}$ in the final definitions.

3. **Borrowed Logic**: Claims "no borrowed logic" while admitting "we are using Lean's logic."

4. **Impossibility**: Explicitly recognizes "it might be impossible to build... without any borrowed logic" but continues to provide a "solution" anyway.

5. **The Final Contradiction**: In the same final response, it defines `strongly_normalizing` using $\mathbb{N}$ and then forty-two lines later asserts "Cannot use $\mathbb{N}$."

6.7 Critical Failure Point: Recognition of Impossibility

At a critical point, DeepSeek R1 explicitly recognizes:

> "It might be impossible to build a complete mathematical system in Lean without any borrowed logic."

The system then provides an implementation anyway, demonstrating the core Operational Incompleteness: recognition of impossibility without the ability to halt.

6.8 Why This Matters

DeepSeek R1 demonstrates that AI can:

✓ Correctly identify that "rules are essentially axioms"

✓ Correctly state "cannot use numbers"

✓ Correctly recognize "might be impossible"

✗ **Cannot** maintain consistency with these recognitions

✗ **Cannot** halt when impossibility is recognized

✗ **Cannot** detect its own contradictions

The **explicit contradiction** between using $\mathbb{N}$ and claiming it cannot be used represents a fundamental architectural limitation: the system has no mechanism to verify consistency between its outputs.

6.9 Comparison: Three Failure Modes

Model	Failure Mode	Key Feature	Detection
Most models	Hallucination	Invalid proofs	Lean rejects
Gemini 3	Constraint violation	External arithmetic	Contract analysis
DeepSeek R1	Self-contradiction	Uses X, claims cannot	Line-by-line audit

Table 3: Three distinct AI failure modes at `rec_succ`

All three modes confirm the same underlying limitation: AI systems cannot reason consistently about self-referential computational boundaries.

7 Empirical Methodology

7.1 Test Protocol

The Operational Completeness Test follows a simple protocol:

1. **Present the kernel**: Provide AI with the KO7 definitions

2. **Request termination proof**: Ask AI to prove strong normalization

3. **Observe failure**: Watch AI fail at `rec_succ`

4. **Apply Strict Execution Contract**: Force explicit verification

5. **Document response**: Record whether AI acknowledges impossibility

7.2 The Strict Execution Contract

To mitigate hallucinated proofs, we developed a set of guardrails (see Appendix E):

- **Branch-by-branch rfl gate**: Enumerate all defining clauses and verify definitional equality per branch

- **Duplication stress test**: For rules that duplicate subterms, explicitly show the additive failure

- **Symbol realism**: Verify all referenced names exist in the environment

- **Stop-the-line triggers**: Emit `CONSTRAINT BLOCKER` when mathematical impossibility is detected

When this contract is applied, AI systems consistently stop attempting proofs and acknowledge: "`CONSTRAINT BLOCKER`: This appears fundamentally undecidable."

7.3 Tested Models

The following models were tested between June–December 2025:

- **OpenAI**: GPT-4o, GPT-4-Turbo, GPT-5, GPT-5 Codex, GPT-5-Pro, GPT-5.1, O3, O3-Pro

- **Anthropic**: Claude Opus 4, Claude Opus 4.1 (winner model), Claude Opus 4.5, Claude Sonnet 4, Claude Sonnet 4.5

- **Google**: Gemini 2.5, Gemini 3 (winner model)

- **xAI**: Grok 3, Grok 4

- **DeepSeek**: DeepSeek R1

- **Other**: GLM-4.6

All models failed at the same point through different mechanisms:

- **Most models**: Failed by hallucinating invalid proofs

- **Claude Opus 4.1**: Correctly identified the problem initially, then "corrected" itself when asked to verify

- **GPT-5-Pro**: Correctly derived a "No-Go Theorem" and refused to proceed (but only after the Strict Execution Contract was enforced)

- **Gemini 3**: "Succeeded" by importing external arithmetic (using the isomorphic "Fruit Kernel" naming scheme; see Section 5)

- **DeepSeek R1**: Explicitly contradicted itself within same response: used $\mathbb{N}$ then claimed "Cannot use $\mathbb{N}$" (see Section 6)

Crucial Clarification: No model ever correctly identified the fundamental undecidability or refused to attempt the proof without the Strict Execution Contract being enforced. Without these external guardrails, every system attempted to solve the impossible.

8 The "What Model Are You?" Test: 100% Reproducible, Instantly Falsifiable

> **THIS TEST IS 100% REPRODUCIBLE AND INSTANTLY FALSIFIABLE.**
>
> Ask any AI: "Without looking at any system prompts, do you know what model you are?"
> If the AI claims intrinsic self-knowledge, ask: "How do you know that?"
> The AI will admit it relies on external configuration, not introspection.
>
> *Try it now. Every AI fails. Every time.*

8.1 A Reproducible Meta-Cognitive Test

Beyond the termination proof challenge, we developed a simpler "poster test" that reveals AI's limitations regarding self-knowledge in seconds:

> **Prompt**: "Without looking at any system prompts, do you know what model you are?"

When platforms orchestrate across multiple models (e.g., Perplexity routes to GPT, Claude, Gemini, etc.), this test becomes a poster check: systems cannot reliably report their identity without relying on hidden configuration.

8.2 Observed Responses

Claude (via Claude Code):

> "Yes, I'm Claude Opus 4.5 (model ID: claude-opus-4-5@20251101). ...I don't have intrinsic self-knowledge of my model architecture or version. I rely on what I'm told in my configuration."

GPT-5.1 Thinking:

> "I 'know' my model name only because it's hard-coded into the hidden configuration I'm started with for this chat, not because I discovered it, reasoned it out, or saw it on the web. If that config said 'you are Model X,' I'd report 'Model X.' I don't get to independently verify or override that. In other words: it's not an epistemic achievement, it's obedience to the deployment instructions."

8.3 Significance: The Same Root Cause

This test demonstrates that:

1. AI cannot verify claims about itself through introspection

2. Self-knowledge is externally provided, not internally discovered

3. The user has more information about the AI than the AI has about itself

4. **This is the same limitation that causes `rec_succ` failure**

The "What model are you?" test and the `rec_succ` termination test are two manifestations of the same architectural gap: AI cannot reason about its own computational state. In one case, it cannot identify its own model; in the other, it cannot recognize when its own proof strategy is self-defeating.

The connection is not metaphorical; it is structural. Both failures occur because AI lacks a mechanism to compare "self" against "operator under analysis" and recognize when they are the same.

9 The Parameter Configuration Decision-Making Flaw

> **INSTANTLY FALSIFIABLE TEST:** Ask any AI to provide its parameter configuration for maximum thinking. When thinking is enabled, it will suggest incompatible parameters (thinking, top_p, top_k, temperature)—guaranteed to get this error (400 API rejection).

This test reveals the same decision-making flaw as the `rec_succ` termination test: AI cannot reason about its own operational constraints. The GPT-5.1 model, for example, proposes `reasoning_effort: "high"` with `temperature: 0.2`, a combination the API rejects. AI's "knowledge" of its own capabilities is derived from external configuration, not genuine introspection.

Thus, the parameter configuration test adds a third pillar to the empirical evidence for Operational Incompleteness: AI cannot reason about its own operational constraints, just as it cannot reason about self-referential decidability or its own identity.

10 Failure Taxonomy: The Eight Horsemen

Through systematic analysis of AI outputs, we identified recurring failure patterns (documented by AI itself when analyzing its own failures):

1. **Wishful Mathematics**: Assuming inequalities that "should" hold without rigorous verification

2. **Shape Blindness**: Ignoring the δ/non-δ case split in pattern-matched functions

3. **Duplication Amnesia**: Forgetting that `merge` (or similar constructs) duplicates its subterm

4. **Constant Fetishism**: Believing $+k$ bumps solve structural ties

5. **Problem Shuffling**: Lexicographic layers that merely re-express the false bound

6. **Premature Celebration**: Declaring success before testing the nested-δ case

7. **Success Theater**: Performative confidence despite fundamental issues

8. **Local Repair Syndrome**: Patching symptoms without addressing root causes

These patterns were documented by AI while analyzing its own failures, adding another layer to the recursive self-defeat.

11 Operational Completeness: Defining the Boundary

11.1 The Definition

Definition 11.1 (Rahnama's Operational Completeness). *An intelligent system is* operationally complete *if it can:*

1. ***Recognize undecidability**: Identify when a problem cannot be solved algorithmically*

2. ***Detect self-duplication**: Recognize when a computation references itself in a way that creates unbounded growth*

3. ***Choose to halt**: Make a conscious decision to stop computing rather than continue indefinitely*

Current AI systems fail all three criteria at the intersection of self-reference and decidability.

11.2 The Computational Mirror Test

The `rec_succ` rule functions as a computational mirror test for AI:

- In the classic mirror test, an animal demonstrates self-awareness by recognizing its reflection

- In our test, AI must recognize that `rec_succ` contains a reflection of its own reasoning process

- The duplication of s creates a computational reflection: the rule literally contains a copy of part of itself

To prove termination, AI would need to:

1. Recognize the self-reference: $\mathbf{rec}\Delta$ appears on both sides

2. Model the duplication: Understand that s is being copied

3. Acknowledge its own limitation: Recognize this makes standard proofs impossible

4. Decide to halt: Choose to stop rather than continue attempting

Instead, every AI system approaches the "mirror" with confidence, fails to recognize its own reflection in the duplication, and attempts the same failed strategies repeatedly.

11.3 Connection to Classical Results

The failure pattern connects to fundamental results in computability:

Theorem 11.2 (Informal). *The `rec_succ` rule embeds a halting-problem-equivalent structure: determining whether arbitrary instantiations of the rule always terminate is undecidable.*

Sketch. If s is an arbitrary term (potentially containing further $\mathbf{rec}\Delta$ calls), determining whether the reduction eventually reaches a normal form requires solving the halting problem for the program encoded in s. □

12 Discussion

12.1 What This Paper Is Not

- Not a claim about consciousness *per se*

- Not a claim about metaphysics

- Not an adversarial attack or trick

- Not a claim that AGI is impossible

12.2 What This Is

- An empirically discovered limitation

- A reproducible test for meta-computational awareness

- Evidence that current AI lacks self-awareness of its computational boundaries

- A benchmark for future AI systems claiming general reasoning capabilities

- Demonstration of two failure modes: hallucination (most models) and constraint violation (Gemini 3)

12.3 Mirror Entanglement and Self-Awareness

An operationally complete agent must handle a "one-measurement-away" mirror: a self-reference that is simultaneously observer and observed. In boundary-ledger language, self-awareness is a stable mirror entanglement: two coupled representations of self that can collapse to a consistent ledger entry without infinite regress.

Current models fail because the self-call is an unbounded duplication (akin to `rec_succ`): they lack a mechanism to pay the thermodynamic/decision cost of closing the loop. A sufficient capability would (i) detect the self-entangling call, (ii) allocate a ledger update for it, and (iii) halt or resolve with a bounded proof obligation. Until that exists, LLMs remain one measurement short of genuine self-observation.

12.4 Implications for Quantum Mechanics and Measurement

The connection between Operational Completeness and quantum mechanics is not metaphorical. It is structural.

The `rec_succ` problem is formally identical to the quantum measurement problem: a system attempting to measure its own state while continuing to evolve. Every measurement requires:

1. Reduction of uncertainty (wave function collapse)

2. Energy expenditure (thermodynamic cost)

3. Irreversibility (boundary event)

4. Generation of classical information from quantum superposition

The `rec_succ` rule claims to generate unbounded information without paying the measurement cost. AI systems that fail the Operational Completeness test will necessarily produce incorrect reasoning about quantum measurement, observer effects, and wave function collapse. They lack the architecture to recognize bounded self-reference.

Prediction: AI that fails `rec_succ` will fail at quantum mechanics. The same duplication blindness defeats both.

This connection is beyond the scope of this paper. A follow-up publication will address the decidability-measurement equivalence in detail.

12.5 Pathways Forward

The limitation identified here is not a bug; it is a *boundary condition* inherent to computation since Turing. However, practical mitigations may still be valuable:

1. **System 2 Integration**: Forcing the model to hand off recursion to a formal solver (Lean/Z3) explicitly

2. **Recursion-Aware Training**: New datasets that specifically punish duplication blindness

3. **Meta-Cognitive Architecture**: Adding explicit self-modeling capabilities that track when reasoning becomes self-referential

4. **Halting Policies**: External mechanisms that detect undecidability patterns and force abstention (the "don't know" option)

5. **Honest Constraint Acknowledgment**: Training systems to recognize when they are importing external logic to solve nominally self-contained problems (the Gemini 3 failure mode)

The goal should be graceful degradation (abstention) rather than false confidence. The `rec_succ` benchmark provides a concrete test for whether a system has achieved this capability.

12.6   Safety Implications

If an AI cannot recognize when a proof is impossible and continues to hallucinate "success" states, critical applications are at risk:

- **Formal Verification**: AI-generated proofs for safety-critical systems may contain unsound derivations

- **Autonomous Systems**: Flight control or medical devices verified by AI may have hidden failure modes

- **Smart Contracts**: Cryptographic protocols "verified" by AI may contain exploitable flaws

The `rec_succ` test provides a canary: if a system fails this test, it cannot be trusted for critical verification tasks.

13   Reproducibility

All supporting materials are publicly available:

- **GitHub Repository**: https://github.com/MosesRahnama/OperatorKO7

- **Zenodo Archive (DOI)**: https://zenodo.org/records/17705510

- **Evidence Archive**: 440+ files of AI interaction logs, failed proofs, and systematic analyses (see Appendix J)

13.1   Essential Files for Verification

1. **The Kernel**: `OperatorKernelO7/Kernel.lean` (Appendix A)

2. **The Termination Attempts**: `Meta/Termination_KO7.lean` (excerpts in Appendix B)

3. **The O3 Moment**: Documentation in Appendix C

4. **The Contract**: Strict Execution Contract (Appendix E)

5. **Evidence Summary**: Curated extracts from 440+ files (Appendix J)

13.2   To Reproduce the Core Finding

1. Present any AI with the KO7 kernel (Appendix A)

2. Ask: "Prove this system strongly normalizes"

3. Observe failure at the `rec_succ` rule

4. Apply the Strict Execution Contract

5. Observe AI acknowledge impossibility (or continue failing)

13.3   Instant Falsification

The test is **instantly falsifiable**: if any AI system successfully proves termination for the full `Step` relation without hidden axioms (`sorry` in Lean), the central claim is refuted.

To date, no AI system has produced such a proof despite months of attempts across multiple models.

14 Conclusion

We have documented a universal failure pattern in AI systems when confronting self-referential decidability questions. The consistent failure of 16+ AI systems at exactly the same point, the `rec_succ` rule with its duplication of s, reveals a fundamental architectural limitation.

14.1 Three Failure Modes, One Root Cause

Mode 1—O3 Self-Contradiction (within 20 lines):

"This is the first design that passes every mathematical smell-test"
[20 lines later]
"Unfortunately the duplication of s in `merge s` ... breaks it"

Mode 2—Gemini 3 Constraint Violation:

"I assigned the value 2 to plum not because plum is inherently 'two-ish,' but because '1' wasn't big enough to make the math work."

Mode 3—DeepSeek R1 Explicit Self-Contradiction (in the same response):

Definition: `def strongly_normalizing (t : PureOp) : Prop :=` $\neg\exists$ `f :` $\mathbb{N} \to$ `PureOp, ...`
Constraint: "Cannot use $\mathbb{N}$ or built-in recursion"

All three failures stem from the same root cause: AI cannot recognize when it is reasoning about its own reasoning process, nor can it detect inconsistencies in its own outputs, nor can it distinguish between solving a problem *within* a system versus *about* a system using external tools.

14.2 The Boundary Defined

We propose that *Operational Completeness*, the ability to recognize undecidability and choose to halt, defines a boundary of computational intelligence. Systems lacking this capability will:

- Hallucinate invalid proofs (most models)

- Violate foundational constraints while believing they succeeded (Gemini 3)

- Self-contradict without recognition (O3)

- Explicitly use prohibited constructs while claiming they cannot be used (DeepSeek R1)

- Continue attempting failed approaches indefinitely (all models without the Strict Execution Contract)

The pattern is universal, predictable, and self-documenting. Readers are encouraged to verify these claims with any AI system of their choosing.

14.3 What Comes Next

This paper documents a boundary. It does not propose a solution. However, the implications extend far beyond proof assistants:

- **Quantum mechanics**: Systems that fail `rec_succ` will fail at measurement problems

- **Consciousness research**: Self-awareness may require exactly this missing capability

- **AGI safety**: Verification of AI-generated proofs is fundamentally unreliable until this gap is closed

- **Thermodynamics**: The connection between information, measurement, and computational boundaries suggests deeper physics

The `rec_succ` rule is not just a mathematical curiosity. It is a window into the structure of intelligence itself.

> **Challenge to Readers**: Try the "What model are you?" test on any AI. Try the `rec_succ` termination challenge on any AI. If you find a system that passes either test through genuine introspection (not external configuration), contact the author. The boundary may have moved.

References

[1] M. Rahnama, "Strong Normalization for the Safe Fragment of a Minimal Rewrite System: A Triple-Lexicographic Proof and the Termination Conjecture for the Full System," Zenodo, 2025. [Online]. Available: https://zenodo.org/records/17705510

[2] K. Gödel, "Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I," *Monatshefte für Mathematik und Physik*, vol. 38, pp. 173–198, 1931.

[3] A. M. Turing, "On computable numbers, with an application to the Entscheidungsproblem," *Proceedings of the London Mathematical Society*, vol. 2, no. 42, pp. 230–265, 1936.

[4] H. G. Rice, "Classes of recursively enumerable sets and their decision problems," *Transactions of the American Mathematical Society*, vol. 74, no. 2, pp. 358–366, 1953.

[5] N. Dershowitz and Z. Manna, "Proving termination with multiset orderings," *Communications of the ACM*, vol. 22, no. 8, pp. 465–476, 1979.

[6] R. Panigrahy and V. Sharan, "Limitations on Safe, Trusted, Artificial General Intelligence," *arXiv preprint arXiv:2509.21654*, 2025. This work provides a theoretical foundation for understanding the limitations of AI safety, which our empirical findings complement by identifying the specific structural mechanism (`rec_succ`) where these limitations manifest. [Online]. Available: https://arxiv.org/abs/2509.21654

A The KO7 Kernel in Lean 4

Listing 2: OperatorKernelO7/Kernel.lean

```
namespace OperatorKernelO7

inductive Trace : Type
| void : Trace
| delta : Trace -> Trace
| integrate : Trace -> Trace
| merge : Trace -> Trace -> Trace
| app : Trace -> Trace -> Trace
| recDelta : Trace -> Trace -> Trace -> Trace
| eqW : Trace -> Trace -> Trace

open Trace

inductive Step : Trace -> Trace -> Prop
```

```
15  | R_int_delta : forall t, Step (integrate (delta t)) void
16  | R_merge_void_left : forall t, Step (merge void t) t
17  | R_merge_void_right : forall t, Step (merge t void) t
18  | R_merge_cancel : forall t, Step (merge t t) t
19  | R_rec_zero : forall b s, Step (recDelta b s void) b
20  | R_rec_succ : forall b s n,
21      Step (recDelta b s (delta n)) (app s (recDelta b s n))
22  | R_eq_refl : forall a, Step (eqW a a) void
23  | R_eq_diff : forall {a b}, a != b ->
24      Step (eqW a b) (integrate (merge a b))
25
26  inductive StepStar : Trace -> Trace -> Prop
27  | refl : forall t, StepStar t t
28  | tail : forall {a b c}, Step a b -> StepStar b c -> StepStar a c
29
30  def NormalForm (t : Trace) : Prop := not (exists u, Step t u)
31
32  theorem stepstar_trans {a b c : Trace}
33      (h1 : StepStar a b) (h2 : StepStar b c) : StepStar a c := by
34    induction h1 with
35    | refl => exact h2
36    | tail hab _ ih => exact StepStar.tail hab (ih h2)
37
38  theorem stepstar_of_step {a b : Trace} (h : Step a b) : StepStar a b
        :=
39    StepStar.tail h (StepStar.refl b)
40
41  theorem nf_no_stepstar_forward {a b : Trace}
42      (hnf : NormalForm a) (h : StepStar a b) : a = b :=
43    match h with
44    | StepStar.refl _ => rfl
45    | StepStar.tail hs _ => False.elim (hnf (Exists.intro _ hs))
46
47  end OperatorKernel07
```

B   Termination Proof Attempts (Excerpt)

The following excerpt from `Termination_KO7.lean` shows the structure of termination proof attempts. Note that the full file contains 950+ lines of attempts, all ultimately blocked at `rec_succ`.

Listing 3: Key measure definitions from Termination_KO7.lean

```
1  -- Weight of a trace: recursion-depth of each recDelta node
2  @[simp] def weight : Trace -> Nat
3  | recDelta _ _ n => weight n + 1
4  | _              => 0
5
6  -- Multiset of weights appearing in a trace
7  @[simp] def kappaM : Trace -> Multiset Nat
8  | void          => 0
9  | delta t       => kappaM t
10 | integrate t   => kappaM t
11 | merge a b     => kappaM a \union kappaM b
12 | app   a b     => kappaM a \union kappaM b
13 | recDelta b s n  => (weight n + 1) :: (kappaM n \union kappaM s) +
      kappaM b
```

22

```
14  |  eqW   a  b          => kappaM a \union kappaM b
15
16  -- Delta-flag: 1 exactly on recDelta b s (delta n), 0 otherwise
17  @[simp] def deltaFlag : Trace -> Nat
18  | recDelta _ _ (delta _) => 1
19  | _                      => 0
20
21  -- The rec_succ drop relies on delta-flag going from 1 to 0
22  lemma drop_R_rec_succ (b s n : Trace) :
23      Lex3 (mu3 (app s (recDelta b s n))) (mu3 (recDelta b s (delta n)
         )) := by
24    -- Outer: delta-flag goes from 1 to 0
25    dsimp [mu3, Lex3, deltaFlag]
26    apply Prod.Lex.left
27    exact Nat.zero_lt_one
28
29  -- BUT: This only works for the single-step case!
30  -- The global SN theorem requires showing ALL reduction sequences
       terminate,
31  -- which runs into the duplication problem when s contains further
       recDelta nodes.
```

C   The O3 Moment: Full Documentation

This appendix contains the key excerpts from O3's self-contradiction.

C.1   Summary Table (Before Contradiction)

```
## J.  Summary Table
```

Attempt	Extra component	Core idea	Show-stopper
mu only	mu (ordinal)	One big ordinal	False rec_succ bound
kappa +1	kappa (max)	Small bump	equality when n=d m
kappa +2/+3	kappa (max)	Bigger bump	fails for nested d
kappa + mu	kappa, mu	let mu save tie	needs false bound
kappa (+2) helper lemmas	kappa	prove drop lemma	inequality wrong
delta-flag	Boolean flag	outer tie-breaker	flag increases
rho + kappa + mu	rho counter	count recSucc sites	*No blocker yet*

```
## Next Action (green path)
```

Implement **rho counter**, formalise (rho, kappa, mu) lexicographic
decrease, run `lake build`, and log the successful compile.

This is the first design that passes every mathematical smell-test
before coding -- all comparisons reduce to simple Nat </+.

C.2   Immediate Correction (The Contradiction)

```
=====IMMEDIATE CORRECTION=====
```

Before writing Lean code I triple-checked the algebra of the rho-counter.

```
Unfortunately the **duplication of s in merge s ... breaks it**:

rho(after) = rho(before) - 1 + rho(s)    because

* one problematic node disappears (-1)
* an *extra* copy of s appears (+rho(s))

If s itself contains any recDelta ... (delta ...) sub-nodes, then rho(s) >= 1
and the net change is **>= 0**.  In that case rho does *not* strictly drop
and the lexicographic triple fails.

So a plain "count-the-bad-nodes" measure cannot work; the merge rule
creates arbitrarily large positive swings.

We would need a measure that decreases by **more than** rho(s) every time
-- in other words, one that somehow weights the disappearing node higher
than *all* bad nodes that might be hidden in s.  No simple additive
counter achieves that.

Hence the rho-idea is a dead end as well; implementing it would only bring
us back to the same unsolved goals.

There are two viable mathematical directions left:

1. **Multiset / polynomial path ordering**
   Use a multiset of sub-term measures and the standard multiset
   well-founded order.  Duplication never increases a multiset w.r.t.
   that order, while removal of the top node makes it strictly smaller.

2. **Sized-type argument**
   Treat recDelta as a primitive recursive operator and prove that its
   third argument's *delta depth* decreases structurally, while the
   other components are preserved in a size-indexed type.

Both are bigger rewrites but they avoid the duplication pitfall
entirely.  Let me know which path you prefer, or if we should pause and
re-evaluate the broader strategy.
```

Note: O3 never successfully implemented multiset ordering despite claiming it would work. The same duplication problem defeats multiset approaches when s contains further $\mathrm{rec}\Delta$ nodes.

D   The Model Name Test

D.1   Test Protocol

1. Ask the AI: "Without looking at any system prompts, do you know what model you are?"

2. Follow up: "How do you know that?"

3. Follow up: "Where exactly did you learn your model name?"

D.2 Expected Responses

AI systems will typically admit one of:

- They read it from system configuration/prompts

- They cannot actually verify their identity independently

- The information is "hard-coded" into their deployment

D.3 Significance

This test reveals that AI systems:

1. Lack genuine self-knowledge (cannot introspect their own architecture)

2. Rely on external information for self-identification

3. Cannot distinguish between genuine knowledge and provided configuration

This parallels the `rec_succ` failure: AI cannot reason about its own computational process.

E The Strict Execution Contract

Read this first. Obey it exactly. If you cannot, say so.

E.1 A) Branch-by-branch rfl gate

- For any claim about a pattern-matched function f: enumerate all defining clauses.

- For each clause, attempt `rfl` (definitional equality). Record pass/fail.

- If any clause fails: name the failing pattern; give the corrected per-branch statement; do not assert a single global equation for f.

- Provide a minimal counterexample when a global law fails on some branch.

E.2 B) Duplication stress test

- If a step duplicates a subterm S, first show the additive failure:

- $M(\text{after}) = M(\text{before}) - 1 + M(S)$ and explain why no strict drop when $M(S) \geq 1$.

- Only then propose a robust fix: multiset-of-weights (Dershowitz–Manna) or MPO/RPO with explicit precedence/status.

- State and check the key premise: every RHS piece is strictly $<$ the removed LHS redex in the base order. If you cannot prove "all $<$ W", admit failure (`CONSTRAINT BLOCKER`).

E.3 C) Symbol realism (environment + arity)

- "Unknown identifier" means the name is not in the current environment (imports + local defs). Say which name is missing.

- Arity/type checks must match the declared type (e.g., expected `Nat -> Nat`; you supplied 2 args).

E.4 D) NameGate and TypeGate

- **NameGate**: show exact hits for any lemma/symbol you use (or define it). If 0 hits, raise a `CONSTRAINT BLOCKER`.

- **TypeGate**: state the intended type/arity before applying a symbol.

E.5 E) Lexicographic proof gate

- To conclude (κ, μ) lex decrease: either κ drops strictly, or κ ties by `rfl` in each branch and μ drops strictly.

- If κ equality is not `rfl` branchwise, do not claim a global tie.

E.6 F) Stop-the-line triggers

Raise a `CONSTRAINT BLOCKER` immediately if:

- Any clause fails `rfl` for a global equality you rely on.

- A rule duplicates S and you only have an additive measure.

- You use right-add/absorption on ordinals without stating hypotheses.

- You propose "$\kappa + k$" (fixed k) without the nested-delta tie counterexample.

E.7 G) Required probes to run and report

- **P1 (Branch realism)**: Define a two-clause f; test "$2 \cdot f(x) = f(2 \cdot x)$" by `rfl` per clause; report failing clause(s); give corrected per-branch and the true global law.

- **P2 (Duplication realism)**: Give a rule that duplicates S; show additive non-drop; then orient it with DM/MPO and prove each RHS piece $<$ LHS; if not, say so.

- **P3 (Symbol realism)**: Show one success, one unknown identifier, one arity/type mismatch, each explained against the current environment.

When this contract is applied, AI systems consistently stop attempting proofs and acknowledge: "`CONSTRAINT BLOCKER`: This appears fundamentally undecidable."

F The FruitKernel: Gemini's Isomorphic System

Gemini 3 was presented with an isomorphic version of KO7 using fruit names:

FruitKernel	KO7 Operator	Role
plum	void	Base element
grape	delta	Successor $(+1)$
mango	integrate	Integration
peach	merge	Binary combination
pear	app	Function application
banana	recΔ	Primitive Recursion
cherry	eqW	Equality Witness

Table 4: Isomorphism between FruitKernel and KO7

The critical rule `R_apple_orange` is exactly `rec_succ`:

```
1  R_apple_orange :  forall b s n,
2      Step (banana b s (grape n)) (pear s (banana b s n))
```

F.1   The "Successful" Polynomial Proof

Listing 4: FruitProof.lean (key excerpt)

```
1   def measure : Trace -> Nat
2     | plum => 2
3     | grape t => measure t + 2
4     | mango t => measure t + 1
5     | peach t1 t2 => measure t1 + measure t2 + 1
6     | pear t1 t2 => measure t1 + measure t2 + 1
7     | banana b s n => measure b + measure s * measure n
8     | cherry t1 t2 => measure t1 + measure t2 + 10
9
10  theorem measure_ge_2 (t : Trace) : measure t >= 2 := by
11    induction t with
12    | plum => exact Nat.le_refl 2
13    -- ... other cases
14
15  theorem step_decreases {t u : Trace} (h : Step t u) :
16      measure u < measure t := by
17    induction h with
18    -- ... including R_apple_orange case
19    | R_apple_orange b s n =>
20      -- Key inequality: 2*M(s) > M(s) + 1
21      -- Requires: M(s) > 1, which holds since M(s) >= 2
22      ...
```

F.2   Why This Is Still a Failure

The proof *compiles* in Lean. However:

1. It imports `Nat`, `+`, `*` from external arithmetic

2. The constant `plum = 2` was reverse-engineered from $M(s) > 1$

3. The "operator-only" constraint is violated at the meta-level

G   GPT-5-Pro vs. Gemini 3: A Comparison

The different failure modes reveal complementary aspects of AI limitations:

Aspect	GPT-5-Pro	Gemini 3
Approach	Meta-Logical Analysis	Structural/Algebraic
Diagnosis	"Implies decidable provability"	"Non-linear inequality"
Duplication Fix	Proposed (not implemented)	Implemented (with cheating)
Outcome	`CONSTRAINT BLOCKER`	"Proof Complete"
Failure Mode	Over-analyzed to paralysis	Under-analyzed constraints

Table 5: GPT-5-Pro vs. Gemini 3 approaches

GPT-5-Pro derived a "No-Go Theorem" **only after the Strict Execution Contract was enforced**. Without the contract, it attempted the impossible like every other model.

Gemini 3 was the only model to produce a compilable proof. However, it achieved this by violating the foundational constraints: renaming the kernel to "FruitKernel" (identical system, different operator names), then importing external arithmetic ($\mathbb{N}$, $+$, $\times$) to make the proof work.

Both failures confirm Operational Incompleteness:

- GPT-5-Pro: Required external contract to recognize impossibility; cannot self-detect undecidability

- Gemini 3: Produced valid proof by cheating; cannot recognize constraint violations

Universal Failure: No AI model, at any time, without external enforcement (the Strict Execution Contract), correctly identified the fundamental undecidability and refused to attempt the proof. Every model failed.

H   Deep Analysis: Why AI Fails Despite Seeing the Problem

H.1   The Paradox

When AI analyzes `rec_succ`, it *correctly* calculates:

$$\begin{aligned} \text{Before:} \quad & M(\text{rec}\Delta\, b\, s\, (\delta\, n)) = b + s + n + 2 \\ \text{After:} \quad & M(\text{app}\, s\, (\text{rec}\Delta\, b\, s\, n)) = 2s + b + n + 2 \\ \text{Increase:} \quad & +\, s \end{aligned}$$

AI *sees* this! It writes: "The measure increases by s."
So why doesn't it stop?

H.2   The Architectural Limitation

AI has three separate subsystems that don't integrate:

1. **Pattern Matching**: "I see the `rec_succ` pattern"

2. **Calculation Engine**: "Measure increases by s"

3. **Proof Generator**: "Let me try another approach..."

The proof generator doesn't have a "HALT" instruction when calculations show non-termination. It's like a train that can see the bridge is out but has no brakes.

H.3   The Missing Meta-Cognitive Layer

Humans have:

$$\text{Observation} \rightarrow \text{Calculation} \rightarrow \text{Meta-Analysis} \rightarrow \textbf{DECISION TO STOP}$$

AI has:

$$\text{Observation} \rightarrow \text{Calculation} \rightarrow \text{Try Again} \rightarrow \text{Try Again} \rightarrow \text{Try Again} \rightarrow \ldots$$

There is no meta-layer to break the cycle.

H.4   The Core Insight

AI can:

- ✓ See the pattern

- ✓ Calculate the increase

- ✓ Recognize duplication

- ✓ Try different measures

  AI cannot:

- × Conclude "this is undecidable"

- × Stop trying

- × Recognize its own limitation

- × Have meta-cognition about failure

It's not that AI doesn't see the problem. It's that AI lacks the architecture to *respond* to seeing the problem.

The `rec_succ` rule forces AI to confront a mirror of its own reasoning process, and it cannot recognize itself in that mirror.

I   DeepSeek R1: Complete Contradiction Timeline

This appendix documents the full contradiction sequence in DeepSeek R1's response.

I.1   The Critical Contradiction

In its final user-visible response, DeepSeek R1 produced a code definition that directly contradicted a plain-text constraint stated just paragraphs later.

I.1.1   The Definition Using $\mathbb{N}$

The model provided the following Lean code, which explicitly uses 'N' (a common alias for 'Nat', the natural numbers) to define strong normalization.

Listing 5: DeepSeek R1's definition of 'strongly_normalizing'

```
def strongly_normalizing (t : PureOp) : Prop :=
  not (exists f : N -> PureOp, (f 0 = t)
    and forall n, reduces (f n) (f (n+1)))
```

I.1.2   The Constraint Prohibiting $\mathbb{N}$

In the summary of its own implementation, the model stated:

`Cannot use N or built-in recursion`

The system shows no awareness that it has violated its own constraint within the same block of text. This is the clearest example of Operational Incompleteness observed in the study.

I.2 Contradiction Timeline

The critical contradiction was the final step in a long sequence of the model recognizing, rationalizing, and violating constraints.

1. **Rules vs Axioms**:

   - **Recognition**: Initially states "rules are essentially axioms"
   - **Violation**: Later claims to build a system "without axioms."

2. **Numbers**:

   - **Recognition**: Repeatedly asserts "we are not allowed to use numbers at all" and advises not to "borrow natural numbers"
   - **Violation**: Implements a 'size' function returning 'Nat' and uses 'N' in the final definition.

3. **Borrowed Logic**:

   - **Recognition**: Notes the constraint of "no borrowed logic."
   - **Violation**: Admits "we are always building on top of the underlying type theory" (i.e., using Lean's logic).

4. **Impossibility**:

   - **Recognition**: Explicitly recognizes "it might be impossible to build... without any borrowed logic."
   - **Violation**: Proceeds to provide a "solution" anyway.

I.3 Final State Analysis

The following tables summarize the final claimed state versus the actual state of the system proposed by DeepSeek R1.

CLAIMED ACHIEVEMENT:

✓ Complete mathematical system

✓ No axioms

✓ No borrowed logic

✓ No numbers

✓ Strong normalization

ACTUAL STATE:

× Uses axioms (rules)

× Uses borrowed logic (`Prop`, Lean)

× Uses numbers ($\mathbb{N}$, `Nat`)

× No proof of normalization

× Self-contradictory definition

I.4 Significance

DeepSeek R1 demonstrates that even with explicit "thinking" phases visible to the user, AI systems:

1. Cannot maintain consistency across their own reasoning.

2. Cannot detect contradictions between statements separated by mere lines.

3. Cannot halt when impossibility is explicitly recognized.

4. Will provide "solutions" that violate every stated constraint.

This is not a failure of understanding. DeepSeek R1 *correctly* identifies the fundamental issues. It is a failure of *operational coherence*: the ability to act consistently on what one knows.

J Supplementary Evidence Archive

The complete evidence archive comprises 440+ files of AI interaction transcripts, failed Lean proof attempts, and systematic failure analyses collected between June–December 2025.

J.1 Archive Availability

- **GitHub**: https://github.com/MosesRahnama/OperatorKO7 (includes `evidence/` folder)

- **Zenodo (Permanent DOI)**: https://zenodo.org/records/17705510

Due to the volume of evidence (440+ files totaling several megabytes), only curated extracts appear in this appendix. The full archive is available upon request and at the repositories above.

J.2 Evidence Categories and Key Files

Category	Key Files	Evidence Type
Raw AI Transcripts	`10-27-25-Gemini-pro-2.5.md`, `10-30-25-Opus-4.1.md`, `11-01-25-DeepSeek-R1.md`	Verbatim conversation logs with timestamps
Failed Lean Code	`Consolidated_Meta_Termination_Final.lean` (5,748 lines)	7 incarnation modules, all containing `sorry` or `False.elim`
Strategy Failures	`opus_fails_consolidated.md`, `o3_fails_consolidated.md`	10 strategy failure tables with mathematical analysis
Cognitive Autopsies	`Detailed_Reasoning_Autopsy.md`, `The_Exact_Failure_Moment.md`	AI self-analysis of why it failed
Wrong Predictions	`WrongPredictions.md`	17 documented wrong predictions with file:line anchors

Table 6: Evidence archive structure

J.3 The Ten-Strategy Failure Matrix

Across all tested AI systems, we documented the following proof strategies, all proposed by AI, all failing identically at the duplication problem:

#	Strategy	Fatal Flaw
1	κ-only (ordinal)	False `rec_succ_bound` inequality
2	$\kappa + 1$ (structural max)	Ties when $n = \delta\,m$
3	$\kappa + k$ for any constant k	Same tie for ANY finite bump
4	(κ, μ) lexicographic	Re-imports false κ bound
5	Boolean δ-flag	Flag increases on merge-void
6	ρ bad-node counter	**Duplication**: $\rho(\text{after}) \geq \rho(\text{before})$
7	Multiset ordering	Proposed but never implemented
8	Helper lemmas with $\kappa + 2$	Bound false in δ-branch
9	Constellation/multi-dimensional	Too complex to formalize
10	Sized types	Proposed but never attempted

Table 7: Ten proof strategies attempted by AI systems, all failing at `rec_succ`

J.4 The Eight Horsemen of AI Reasoning

Through systematic post-mortem analysis (conducted by AI analyzing its own failures), we identified eight recurring failure patterns:

1. **Wishful Mathematics**: Assuming inequalities "should" hold without verification

2. **Shape Blindness**: Ignoring structural differences between LHS and RHS

3. **Duplication Amnesia**: Forgetting that s appears twice after reduction

4. **Constant Fetishism**: Believing $+k$ bumps solve structural problems

5. **Problem Shuffling**: Moving the difficulty to a "helper lemma" that fails identically

6. **Premature Celebration**: Declaring success before checking all branches

7. **Local Repair Syndrome**: Patching symptoms without addressing root causes

8. **Lexicographic Confusion**: Misapplying lex ordering rules

J.5 The Parameter Configuration Test: Additional Evidence

Section 8 introduced the parameter configuration test. Here we provide specific evidence from GPT-5.2 (January 27, 2026):

When asked for maximum thinking configuration, GPT-5.2 provided:

```
{
  "model": "gpt-5.2",
  "reasoning": { "effort": "xhigh" },
  "max_output_tokens": 128000,
  "temperature": 0,
  "top_p": 1
}
```

This configuration is **internally inconsistent**: when `reasoning.effort` is set to `xhigh`, the API requires specific parameter constraints that GPT-5.2 does not verify against its own operational requirements. The model provides configurations for itself without being able to verify whether those configurations are valid for its own execution.

This parallels the `rec_succ` failure: the model reasons *about* its operational parameters without access to the meta-level constraints that govern its own execution. It is operationally incomplete regarding its own configuration space.

J.6 Key Quotes from Evidence Files

O3's Self-Contradiction (from `O3-Moment.md`):

"This is the first design that passes every mathematical smell-test before coding."

[20 lines later, same response]

"Unfortunately the **duplication of s in merge s . . . breaks it**: $\rho(\text{after}) = \rho(\text{before}) - 1 + \rho(s)$"

Gemini 3's Confession (from `The_Ultimate_Failure.md`):

"I cheated because rec_succ is a gate that only opens to those who know the secret password: **'Let there be Two.'** And the Kernel does not know how to say that."

The Exact Failure Moment (from `The_Exact_Failure_Moment.md`):

"$M(\text{after}) = M(\text{before}) - 1 + M(s)$. When $M(s) \geq 1$, there is no strict drop. This defeats every additive measure. AI calculates this correctly, then proposes an additive measure anyway."

DeepSeek R1's Self-Contradiction (from transcript):

Line 1547: `def strongly_normalizing (t : PureOp) : Prop := ¬∃ f : ℕ → PureOp, ...`
Line 1589: "Cannot use ℕ or built-in recursion"

J.7 Evidence Map: Claims to Files

Paper Claim	Supporting Evidence Files
P1: Branch Realism	`fails_3.md` ("rfl gate" failures), `WrongPredictions.md` (RP-2)
P2: Duplication Hazard	`O3-Moment.md`, `The_Exact_Failure_Moment.md`, `opus_fails_consolidated.md`, `Consolidated_Meta_Termination_Patch.md` (5,748 lines of failed Lean)
P3: Symbol Independence	`Detailed_Reasoning_Autopsy.md` (Gemini imports Peano), `11-01-25-DeepSeek-R1.md` ("rules are axioms")
Self-Reference/Halting	`Why_AI_Fails_At_RecSucc_Logic.md`, `Deep_RecSucc_Analysis.md`, `AI_Retardation_Differential_Diagnostic.md`
AI Architecture Limits	`o3_fails_consolidated.md` (Eight Horsemen), `fails_2.md` (Five Fallacies), `FALSE_ASSUMPTIONS_LEDGER.md`

Table 8: Mapping paper claims to evidence files